

# Unsupervised Learning for Solving the Travelling Salesman Problem

Yimeng Min<sup>\*1</sup> Yiwei Bai<sup>\*1</sup> Carla P. Gomes<sup>1</sup>

## Abstract

We propose UTSP, an unsupervised learning (UL) framework for solving the Travelling Salesman Problem (TSP). We train a Graph Neural Network (GNN) using a surrogate loss. The GNN outputs a heat map representing the probability for each edge to be part of the optimal path. We then apply local search to generate our final prediction based on the heat map. Our loss function consists of two parts: one pushes the model to find the shortest path and the other serves as a surrogate for the constraint that the route should form a Hamiltonian Cycle. Experimental results show that UTSP outperforms the existing data-driven TSP heuristics. Our approach is parameter efficient as well as data efficient: the model takes  $\sim 10\%$  of the number of parameters and  $\sim 0.2\%$  of training samples compared with reinforcement learning or supervised learning methods.

## 1. Introduction

Euclidean Travelling Salesman Problem (TSP) is one of the most intensely studied NP-hard graph problems in combinatorial optimization community. Traditional methods, such as Concorde (Applegate et al., 2006) and LKH (Helsgaun, 2000), use linear programming and search to solve TSP. Recently, the success of Graph Neural Networks (GNNs) for a variety of machine learning tasks has sparked interest in building data-driven heuristics for approximating TSP solutions with lower time costs. These learning-based models usually build heuristics by reducing the length of TSP tours via reinforcement learning (RL) or directly learning from the optimal solutions via supervised learning (SL). However, since TSP is NP-hard, SL can cause expensive annotation problems due to the costly search times involved in generating optimal solutions. While for RL, when dealing with big graphs, the model will run into the sparse reward problem because the reward is only decided after decoding

a complete solution. The sparse reward problem results in poor generalization performance and high training variance. Furthermore, both RL and SL suffer from costly large-scale training. These models take more than one million training samples when dealing with TSP with 100 nodes, making the training process very time-consuming (Kool et al., 2019)(Joshi et al., 2022)(Qiu et al., 2022).

In this work, we build a data-driven TSP heuristic in an unsupervised learning fashion. We construct a surrogate loss function with two parts: one encourages the GNN to find the shortest path and the other acts as a proxy for the constraint that the path should be a Hamiltonian Cycle over all nodes. The surrogate loss enables us to update the model without having to decode a complete solution. This helps alleviate the sparse reward problem encountered in RL and thus it avoids unstable training or slow convergence (Kool et al., 2019). Our UTSP method does not rely on labeled data. This helps the model avoid the expensive annotation problems encountered in SL and significantly reduces the time cost. In fact, due to the prohibitive time cost of building training datasets for large instances, many SL methods are trained on relatively small instances only (Fu et al., 2021)(Joshi et al., 2019a). Such SL models scale poorly to the big instances, while with our UTSP pipeline, we are able to train our model on larger instances directly. Overall, our training does not rely on any labeled training data and converges faster comparing to RL/SL methods.




Figure 1. Illustration of training pipeline. We use a GNN to learn a transition matrix  $\mathbb{T}$  and we build a heat map  $\mathcal{H}$  based on  $\mathbb{T}$ . We combine the transition matrix, heat map and the distance matrix to build the surrogate loss.

Our method is shown in Figure 1. The model takes the coordinates as the input of GNN. The distance between

<sup>\*</sup>Equal contribution <sup>1</sup>Department of Computer Science, Cornell University, Ithaca, NY, 14850, USA. Correspondence to: Yimeng Min <min@cs.cornell.edu>.

two nodes determines the edge weight in the adjacency matrix. After training the GNN, the heat map is converted to a valid tour using local search. We evaluate the performance of UTSP through comparisons on TSP cases of fixed graph sizes up to 1,000 nodes. We note that UTSP is fundamentally different from RL, which may also be considered unsupervised. While RL requires a Markov Decision Process (MDP) and its reward is extracted after obtaining solutions, does not use an MDP and the loss function (reward) is determined based on a heat map. Furthermore, UTSP incorporates a data-driven approach, which learns hidden structure across instances in an unsupervised way, in contrast to standard direct optimization that typically considers one instance at a time, without a learning component across instances. Overall, UTSP requires only a small amount of (unlabeled) data and compensates for it by employing an unsupervised surrogate loss function and an expressive GNN. UTSP outperforms other learning-based methods in terms of performance accuracy and speed. The heat maps built using UTSP help reduce the search space and generate "algorithmic priors" which facilitate the local search. We further show that the expressive power of GNNs is critical for generating non-smooth heat maps.

## 2. Model

In this paper, we study symmetric TSP on 2D plane. Given  $n$  cities and the coordinates  $(x_i, y_i) \in \mathbb{R}^2$  of these cities, our goal is to find the shortest possible route that visits each city exactly once and returns to the origin city, where  $i \in \{1, 2, 3, \dots, n\}$  is the index of the city.

### 2.1. Graph Neural Network

Given a TSP instance, let  $\mathbf{D}_{i,j}$  denote the Euclidean distance between city  $i$  and city  $j$ .  $\mathbf{D} \in \mathbb{R}^{n \times n}$  is the distance matrix. We first build adjacency matrix  $\mathbf{W} \in \mathbb{R}^{n \times n}$  with  $\mathbf{W}_{i,j} = e^{-\mathbf{D}_{i,j}/\tau}$  and node feature  $\mathbf{F} \in \mathbb{R}^{n \times 2}$  based on the input coordinates, where  $\mathbf{F}_i = (x_i, y_i)$  and  $\tau$  is the temperature. The node feature matrix  $\mathbf{F}$  and the weight matrix  $\mathbf{W}$  are then fed into a GNN to generate a transition matrix  $\mathbb{T} \in \mathbb{R}^{n \times n}$ .

In our model, we use Scattering Attention GNN (SAG), SAG has both low-pass and band-pass filters and can build adaptive representations by implicitly learning node-wise weights for combining multiple different channels in the network using attention-based architecture. Recent studies show that SAG can output expressive representations for graph combinatorial problems such as maximum clique while remaining lightweight (Min et al., 2022).

Let  $\mathcal{S} \in \mathbb{R}^{n \times n}$  denote the output of SAG, we first apply a column-wise Softmax activation to the GNN's output and we can summarize this operation in matrix notation as  $\mathbb{T}_{i,j} = e^{\mathcal{S}_{i,j}} / \sum_{k=1}^n e^{\mathcal{S}_{k,j}}$ . This ensures that each element in

$\mathbb{T}$  is greater than zero and the summation of each column is 1. We then use  $\mathbb{T}$  to build a heat map  $\mathcal{H}$ , where  $\mathcal{H} \in \mathbb{R}^{n \times n}$ . In our model, we use  $\mathcal{H}$  to estimate the probability of each edge belonging to the optimal solution and use  $\mathbb{T}$  to build a surrogate loss of the Hamiltonian Cycle constraint. As illustrated in Figure 2, our approach aims to generate an expressive transition matrix  $\mathbb{T}$  which assigns large weights (close to 1) on the transition elements and small weights (close to 0) on others. This will allow us to build a non-smooth heat map  $\mathcal{H}$  and improve the performance of the local search.




Figure 2. We use a SAG to generate a non-smooth transition matrix  $\mathbb{T}$ . The SAG model is a function of the coordinates and the weighted adjacency matrix.

### 2.2. Building the Heat Map using the Transition Matrix

We build the heat map  $\mathcal{H}$  based on  $\mathbb{T}$ . As mentioned,  $\mathcal{H}_{i,j}$  is the probability for edge  $(i,j)$  to belong to the optimal TSP solution or not. We define  $\mathcal{H}$  as:

$$\mathcal{H} = \mathbb{T} \mathbb{V} \mathbb{T}^T, \quad (1)$$

where

$$\mathbb{V} = \begin{pmatrix} 0 & 1 & 0 & 0 & \cdots & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & \cdots & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & \cdots & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & \cdots & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 \end{pmatrix}$$

is the Sylvester shift matrix (Sylvester, 1909),  $\mathbb{V} \in \mathbb{R}^{n \times n}$ . We can interpret  $\mathbb{V}$  as a cyclic permutation operator that performs a circular shift. The elements in  $\mathcal{H}$  can be written as:  $\mathcal{H}_{i,j} = \sum_{k=1}^n \mathbb{T}_{i,k} \mathbb{T}_{j,k+1(\text{mod } n)}$ .  $\mathcal{H}_{i,j}$  is the sum of element-wise multiplication on two terms: the  $i$ -th row of the  $\mathbb{T}$  and the  $j$ -th row of the  $\mathbb{T}$  with left translation. For example, given the transition matrix in Figure 2, Figure 3 illustrates the corresponding heat map and the TSP solution<sup>1</sup>.

<sup>1</sup>We build  $\mathcal{H}$  using Equation 1 instead of directly using  $\mathbb{T}$  because, under binary assumption, the output of SAG does not satisfy the Hamiltonian Cycle constraint. An example is shown in Figure 2.




Figure 3. Left: the heat map corresponds to the transition matrix in Figure 2; right: the corresponding TSP path.

The first row in  $\mathcal{H}$  is the probability distribution of directed edges start from city 1, and since the third element is the only non-zero one in the first row, we then add directional edge  $1 \rightarrow 3$  to our TSP solution. Similarly, the first column in  $\mathcal{H}$  can be regarded as the probability distribution of directed edges which end in city 1. Ideally, given a graph  $\mathcal{G}$  with  $n$  nodes, we want to build a transition matrix where each row and column are assigned with one value 1 (True) and  $n - 1$  values 0 (False), so that the heat map will only contain one valid solution. In practice, we will build a transition matrix  $\mathbb{T}$  whose heat map  $\mathcal{H}$  assigns large probabilities to the edges in the TSP solution and small probabilities to the other edges.




Figure 4. Illustration of building a cycle from transition matrix  $\mathbb{T}$ . Here  $q_{k-1} = 3$ ,  $q_k = n - 2$ ,  $q_{k+1} = 2$ ,  $q_n = n$  and  $q_1 = 4$ , this means  $\mathcal{H}$  contains the following four directed edges:  $3 \rightarrow n - 2$ ,  $n - 2 \rightarrow 2$ ,  $1 \rightarrow n$  and  $n \rightarrow 4$ .

**Lemma 2.1.** Let  $q_i$  denote the row index of the non-zero element in  $i$ -th column in  $\mathbb{T}$ ,  $\mathbb{T}_{q_i, i} = 1$ ,  $q_i \in \{1, 2, 3, 4, \dots, n\}$ . When each row and column in  $\mathbb{T}$  have one value 1 (True) and  $n - 1$  value 0 (False), then  $q_i = q_j$  if and only if  $i = j$ .

*Proof.* If there exist a  $(i, j)$  pair where  $q_i = q_j$  when  $i \neq j$ , then  $\mathbb{T}_{q_i, i} = 1$  and  $\mathbb{T}_{q_j, j} = \mathbb{T}_{q_i, j} = 1$ . This means  $q_i$ -th row has two non-zero elements, which leads to a contradiction.  $\square$

**Lemma 2.2.** Consider graph  $\mathcal{G}$  with  $n$  nodes, for any  $\mathbb{T} \in \mathbb{R}^{n \times n}$  with  $\mathbb{T}_{i, j} \geq 0$ , when each row and column in  $\mathbb{T}$  have one value 1 (True) and  $n - 1$  value 0 (False). Then, each row and column in  $\mathcal{H}$  also have one value 1 (True) and  $n - 1$  value 0 (False), which means each city corresponds to only one beginning point and one ending point.

*Proof.* First, it is clear that for  $\forall a, b$ ,  $\mathcal{H}_{a, b} \in \mathbb{Z}_{\geq 0}$ . Let's assume  $\mathbb{T}_{a, l} = \mathbb{T}_{b, m} = 1 (a \neq b, l \neq m)$ , since  $\mathcal{H}_{i, j} = \sum_{k=1}^n \mathbb{T}_{i, k} \mathbb{T}_{j, k+1(\text{mod } n)}$ , we then have

$$\begin{aligned} \sum_{j=1}^n \mathcal{H}_{a, j} &= \sum_{j=1}^n \sum_{k=1}^n \mathbb{T}_{a, k} \mathbb{T}_{j, k+1(\text{mod } n)} \\ &= \sum_{k=1}^n \mathbb{T}_{a, k} \left\{ \sum_{j=1}^n \mathbb{T}_{j, k+1(\text{mod } n)} \right\} \\ &= \sum_{k=1}^n \mathbb{T}_{a, k} = 1. \end{aligned}$$

This implies that the summation of each row in  $\mathcal{H}$  is 1. Similarly,

$$\begin{aligned} \sum_{i=1}^n \mathcal{H}_{i, b} &= \sum_{i=1}^n \sum_{k=1}^n \mathbb{T}_{i, k} \mathbb{T}_{b, k+1(\text{mod } n)} \\ &= \sum_{k=1}^n \mathbb{T}_{b, k+1(\text{mod } n)} \left\{ \sum_{i=1}^n \mathbb{T}_{i, k} \right\} \\ &= \sum_{k=1}^n \mathbb{T}_{b, k+1(\text{mod } n)} = 1. \end{aligned}$$

This suggests the summation of each column in  $\mathcal{H}$  is 1. Since each element in  $\mathcal{H}_{i, j} \in \mathbb{Z}_{\geq 0}$ , we can then conclude that each row and column in  $\mathcal{H}$  have one value 1 (True) and  $n - 1$  value 0 (False). Also,  $\mathcal{H}_{i, i} = \sum_{k=1}^n \sum_{l=1}^n \mathbb{T}_{i, k+1(\text{mod } n)} \mathbb{T}_{i, k} = 0$ , this indicates that the elements in the main diagonal of  $\mathcal{H}$  are 0, which implies no self-loops. As mentioned, the  $i$ -th row in  $\mathcal{H}$  is the probability distribution of directed edges start from city  $i$ , and the  $j$ -th column is the probability distribution of directed edges end in city  $j$ . Because each row and column in  $\mathcal{H}$  have one value 1 (True) element and  $\mathcal{H}$ 's diagonal entries are all zero, this means that each city is the beginning point of one directed edge and is also the ending point of another different directed edge.  $\square$

**Lemma 2.3.** There is at least one cycle in  $\mathcal{H}$  which contains  $n$  edges and visits all cities when each row and column in  $\mathbb{T}$  have one value 1 (True) and  $n - 1$  value 0 (False).

*Proof.* Lemma 2.1 indicates that  $q_i \neq q_j$  when  $i \neq j$  and  $q_i \in \{1, 2, 3, 4, \dots, n\}$ , then  $\cup_{i=1}^n q_i = \{1, 2, 3, 4, \dots, n\}$ .

From Equation 1, we have

$$\begin{aligned}\mathcal{H}_{q_i, q_{i+1}} &= \sum_{k=1}^n \mathbb{T}_{q_i, k} \mathbb{T}_{q_{i+1}, k+1(\text{mod } n)} \\ &\geq \mathbb{T}_{q_i, i} \mathbb{T}_{q_{i+1}, i+1(\text{mod } n)} \\ &\geq 1.\end{aligned}$$

Using **Lemma 2.2**, since each row and column in  $\mathcal{H}$  have one value 1 (True) and  $n - 1$  value 0 (False), it suffices to show that  $1 \geq \mathcal{H}_{q_i, q_{i+1}} \geq 1$ , therefore  $\mathcal{H}_{q_i, q_{i+1}} = 1$ . This suggests that there is a directed edge from city  $q_i$  to  $q_{i+1}$ . We can then construct a cycle  $\mathcal{C}$  from  $\mathcal{H}$ , we can write  $\mathcal{C}$  as

$$q_1 \rightarrow q_2 \rightarrow q_3 \rightarrow q_4 \rightarrow \dots \rightarrow q_n \rightarrow q_1,$$

where  $\rightarrow$  is a directed edge. Since  $\cup_{i=1}^n q_i = \{1, 2, 3, 4, \dots, n\}$ , cycle  $\mathcal{C}$  visits all  $n$  cities and have  $n$  edges. One example of how to build  $\mathcal{H}_{q_i, q_{i+1}}$  from  $\mathbb{T}$  is shown in Figure 4.  $\square$

**Corollary 2.4.**  $\mathcal{H}$  represents one Hamiltonian Cycle when each row and column in  $\mathbb{T}$  have one value 1 (True) and  $n - 1$  value 0 (False).

*Proof.* From **Lemma 2.3**, cycle  $\mathcal{C}$  contains  $n$  edges and visits all cities, if there exists another edge  $(i, j)$  which does not belong to  $\mathcal{C}$ , then city  $i$  is the starting point of at least two edges and city  $j$  is the ending point of at least two edges. This results in a contradiction with **Lemma 2.2**. Thus, it suffices to conclude that  $\mathcal{C}$  visits each city exactly once and  $\mathcal{H}$  only contains the edges in  $\mathcal{C}$ . This implies that  $\mathcal{H}$  represents one Hamiltonian Cycle.  $\square$

### 2.3. Unsupervised Loss

In order to generate such an expressive transition matrix  $\mathbb{T}$ , we minimize the following objective function:

$$\begin{aligned}\mathcal{L} = & \lambda_1 \underbrace{\sum_{i=1}^n \left( \sum_{j=1}^n \mathbb{T}_{i,j} - 1 \right)^2}_{\text{Row-wise constraint}} + \lambda_2 \underbrace{\sum_i \mathcal{H}_{i,i}}_{\text{No self-loops}} \\ & + \underbrace{\sum_{i=1}^n \sum_{j=1}^n \mathbf{D}_{i,j} \mathcal{H}_{i,j}}_{\text{Minimize the distance}}.\end{aligned}\quad (2)$$

The first term in  $\mathcal{L}$  encourages the summation of each row in  $\mathbb{T}$  to be close to 1. As mentioned, we normalize each column of  $\mathbb{T}$  using Softmax activation. So when the first term is minimized to zero, each row and column in  $\mathbb{T}$  are normalized. The second term penalizes the weight on the main diagonal of  $\mathcal{H}$ , this discourages self-loops in TSP solutions. The third term can be regarded as the expectation TSP length of the heat map  $\mathcal{H}$ , where  $\mathbf{D}_{i,j}$  is the distance between city  $i$  and  $j$ .

As mentioned, since  $\mathcal{H}$  corresponds to one Hamiltonian Cycle given an ideal transition matrix with one value 1 (True) and  $n - 1$  value 0 (False) in each row and column. Then the minimum value of  $\sum_{i=1}^n \sum_{j=1}^n \mathbf{D}_{i,j} \mathcal{H}_{i,j}$  is the shortest Hamiltonian Cycle on the graph, which corresponds to the optimal solution of TSP.

Given a heat map  $\mathcal{H}$ , we consider  $M$  largest elements in each row (without diagonal elements) and set other  $n - M$  elements as 0. Let  $\tilde{\mathcal{H}}$  denote the new heat map, we then symmetrize the new heat map by  $\mathcal{H}' = \tilde{\mathcal{H}} + \tilde{\mathcal{H}}^T$ . Let  $\mathbf{E}_{ij} \in \{0, 1\}$  denote whether an undirected edge  $(i, j)$  is in our prediction or not. Without loss of generality, we can assume  $0 < i < j \leq n$  and define  $\mathbf{E}_{ij}$  as :

$$\mathbf{E}_{ij} = \begin{cases} 1, & \text{if } \mathcal{H}'_{ij} = \mathcal{H}'_{ji} > 0 \\ 0, & \text{otherwise} \end{cases}.$$

Let  $\Pi$  denote the set of undirected edges  $(i, j)$  with  $\mathbf{E}_{ij} = 1$ . Ideally, we would build a prediction edge set  $\Pi$  with a small  $M$  value, and  $\Pi$  can cover all the ground truth edges so that we are able to reduce search space size from  $n(n - 1)/2$  to  $|\Pi|$ . In practice, generating small-size edge sets  $\Pi$  which always cover the ground truth solutions is very difficult. We aim to let  $\Pi$  cover as many ground truth edges as possible and use  $\mathcal{H}'$  to guide the local search process.

## 3. Local Search

### 3.1. Heat Map Guided Best-first Local Search

We employ the best-first local search guided by the heat map to generate the final solution. Best-first search is a heuristic search that explores the search space by expanding the most promising node selected w.r.t. an evaluation function  $f(\text{node})$ . In our framework, each node of the search tree is a complete TSP solution. For the initialization of one search tree, we randomly generate a valid TSP solution and improve it using the 2-opt heuristic until no better solution is found. The expand action of the search node refers to (Fu et al., 2021) and is based on the widely used  $k$ -opt heuristic (Croes, 1958), where it replaces  $k$  old edges (in the current solution) with  $k$  new edges, i.e., transforms the old solution to a new solution. More formally, we use a series of cities,  $u_1, v_1, u_2, \dots, u_k, v_{k+1}$ , to represent an action, where  $v_{k+1} = u_1$  to ensure it is a valid solution. All the edges  $(u_i, v_i)$  ( $1 \leq i \leq k$ ) are removed from the tour and  $(v_i, u_{i+1})$  ( $1 \leq i \leq k$ ) are added to the tour. Note that once we know  $u_i, v_i$  is deterministically decided.  $u_1$  is randomly selected for each expansion, and  $v_1$  is decided subsequently. Then we select  $u_{i+1}$  ( $i \geq 1$ ) as follows: (1) if  $u_{i+1} = u_1$ , i.e., forms a new TSP tour, leads to an improved solution then we set  $u_{i+1} = u_1$  and have a candidate solution. (2) if  $i \geq K$ , then we will discard this action and start a new expand action, where  $K$  is a hyper-parameter

which controls the maximal edges we can remove in one action. (3) otherwise, we select  $u_{i+1}$  based on the **heat map** stochastically. We use  $N_{u,v}$  to denote the times the edge  $(u, v)$  is selected during the entire search procedure. The likelihood of selecting the edge  $(u, v)$  is denoted by  $L_{u,v}$ :

$$L_{u,v} = \mathcal{H}'_{u,v} + \alpha \sqrt{\frac{\log(S+1)}{N_{u,v}+1}},$$

where  $\alpha$  is a hyper-parameter and  $S$  is the local search's total number of expand actions. The first term encourages the algorithm to select the edge with a high heat map value, while the second term diversifies the selected edges. Moreover, when selecting the city  $v$  given  $u$ , we only consider the cities from the candidate set of  $v$ . This candidate set consists of cities with the top  $M$  heat map value or the nearest  $M$  cities

Among all the possible new solutions, we use the tour's length as the evaluation function  $f$ , i.e., we select the solution of the shortest tour length as the next search node. For each search node, we try at most  $T$  expand actions. From these  $T$  expand actions, if no improved solution is found, we randomly generate a new initial solution and start another round of best-first local search.

### 3.2. Updating the Heat Map

We borrow the idea of the backpropagation used in Monte Carlo Tree Search (MCTS). We use  $s$  to denote the current search node,  $s'$  to denote the next search node ( $s'$  has to improve  $s$ ), and  $L(s)$  to represent the tour length of node  $s$ . The heat map  $\mathcal{H}'$  is updated as following:

$$\mathcal{H}'_{v_i, u_{i+1}} = \mathcal{H}'_{v_i, u_{i+1}} + \beta [\exp(\frac{L(s) - L(s')}{L(s)}) - 1],$$

where  $\beta$  is a search parameters and  $(v_i, u_{i+1}) (1 \leq i \leq k)$  is the actions used to transform  $s$  to  $s'$ . We intend to raise the importance of the edges that lead to a better solution. If we cannot find an improved solution for the current node, then no update is executed for the heat map.

### 3.3. Leveraging Randomness

Randomness is shown to be very powerful in the local search community (Gomes et al., 1998; Bresina, 1996; Gomes et al., 2000). Our local search procedure also employs it to improve performance. When we stop the current round of local search by not finding an improved solution within  $T$  expand actions and switch to a new best-first local search with a new initial solution, we randomly modify the parameter  $K$ . A greater value of  $K$  results in spending more time searching from one initial solution. The intuition is that sometimes we want more initial solutions while sometimes we want to search deeper (replace more edges in  $k$ -opt) for a specific

solution. Besides that, we also randomly decide how we construct the candidate set (based on the heat map or pairwise distance) for each city before the new round of local search.

## 4. Experiments

### 4.1. Dataset

Our dataset contains 2,000 samples for training and 1,000 samples for validation. We use the same test dataset in (Fu et al., 2021). The test dataset contains 10,000 2D-Euclidean TSP instances for  $n = 20, 50, 100, 128$  instances for  $n = 200, 500, 1,000$ . We train our models on TSP instances with 20, 50, 100, 200, 500, and 1,000 vertices. We then build the corresponding heat maps based on these trained models.

### 4.2. Results




Figure 5. TSP 100 training curve using unsupervised learning surrogate loss. We compare two GNN models: GCN (Kipf & Welling, 2016) and SAG (Min et al., 2022), where GCN is a low-pass model and SAG is a low-pass + band-pass model.

Table 1 and Table 2 present model's performance on TSP 20, 50, 100, 200, 500 and 1,000. The first three lines in Table 1 and Table 2 summarize the performance of two exact solvers (Concorde and Gurobi) and LKH3 heuristic (Helsgaun, 2017). The learning-based methods can be divided into RL sub-category and SL sub-category. Greedy decoding (G), Sampling (S), Beam Search (BS), and Monte Carlo Tree Search are the decoding schemes used in RL/SL. The 2-OPT is a greedy local search heuristic.

We compare our model with existing solvers as well as different learning-based algorithms. The performance of our method is averaged of four runs with different random seeds. The running time for our method is divided into two parts: the inference time (building the heat map  $\mathcal{H}$ ) and the search time (running search algorithm). On small instances, our results match the ground-truth solutions and generate average gaps of **-0.00009%**, **-0.002%** and **-0.00011%** respectively on instances with  $n = 20, 50, 100$ , where the negative



Table 1. Results of SAG + Local Search w.r.t. existing baselines, tested on 10,000 instances with  $n = 20, 50$  and  $100$ .

Method	Type	TSP20			TSP50			TSP100		
		Length	Gap (%)	Time	Length	Gap (%)	Time	Length	Gap (%)	Time
Concorde	Solver	3.8303	0.0000	2.31m	5.6906	0.0000	13.68m	7.7609	0.0000	1.04h
Gurobi	Solver	3.8302	-0.0001	2.33m	5.6905	0.0000	26.20m	7.7609	0.0000	3.57h
LKH3	Heuristic	3.8303	0.0000	20.96m	5.6906	0.0008	26.65m	7.7611	0.0026	49.96m
GAT (Deudon et al., 2018)	RL, S	3.8741	1.1443	10.30m	6.1085	7.3438	19.52m	8.8372	13.8679	47.78m
GAT (Deudon et al., 2018)	RL, S 2-OPT	3.8501	0.5178	15.62m	5.8941	3.5759	27.81m	8.2449	6.2365	4.95h
GAT (Kool et al., 2019)	RL, S	3.8322	0.0501	16.47m	5.7185	0.4912	22.85m	7.9735	2.7391	1.23h
GAT (Kool et al., 2019)	RL, G	3.8413	0.2867	6.03s	5.7849	1.6568	34.92s	8.1008	4.3791	1.83m
GAT (Kool et al., 2019)	RL, BS	3.8304	0.0022	15.01m	5.7070	0.2892	25.58m	7.9536	2.4829	1.68h
GCN (Joshi et al., 2019a)	SL, G	3.8552	0.6509	19.41s	5.8932	3.5608	2.00m	8.4128	8.3995	11.08m
GCN (Joshi et al., 2019a)	SL, BS	3.8347	0.1158	21.35m	5.7071	0.2905	35.13m	7.8763	1.4828	31.80m
GCN (Joshi et al., 2019a)	SL, BS*	3.8305	0.0075	22.18m	5.6920	0.0251	37.56m	7.8719	1.4299	1.20h
Att-GCRN(Fu et al., 2021)	SL+RL MCTS	3.8300	<b>-0.0074</b>	23.33s + 1.05m	5.6908	0.0032	2.59m + 2.63m	7.7616	0.0096	3.94m + 5.25m
<b>UTSP (ours)</b>	UL, Search	3.8303	-0.0009	38.23s + 1.04m	5.6894	<b>-0.0200</b>	1.34m + 2.60m	7.7608	<b>-0.0011</b>	5.68m+ 5.21m

 Table 2. Results of SAG + Local Search w.r.t. existing baselines, tested on 128 instances with  $n = 200, 500$  and  $1000$ .

Method	Type	TSP200			TSP500			TSP1000		
		Length	Gap (%)	Time	Length	Gap (%)	Time	Length	Gap (%)	Time
Concorde	Solver	10.7191	0.0000	3.44m	16.5458	0.0000	37.66m	23.1182	0.0000	6.65h
Gurobi	Solver	10.7036	-0.1446	40.49m	16.5171	-0.1733	45.63h	-	-	-
LKH3	Heuristic	10.7195	0.0040	2.01m	16.5463	0.0029	11.41m	23.1190	0.0036	38.09m
GAT (Deudon et al., 2018)	RL, S	13.1746	22.9079	4.84m	28.6291	73.0293	20.18m	50.3018	117.5860	37.07m
GAT (Deudon et al., 2018)	RL, S 2-OPT	11.6104	8.3159	9.59m	23.7546	43.5687	57.76m	47.7291	106.4575	5.39h
GAT (Kool et al., 2019)	RL, S	11.4497	6.8160	4.49m	22.6409	36.8382	15.64m	42.8036	85.1519	63.97m
GAT (Kool et al., 2019)	RL, G	11.6096	8.3081	5.03s	20.0188	20.9902	1.51m	31.1526	34.7539	3.18m
GAT (Kool et al., 2019)	RL, BS	11.3769	6.1364	5.77m	19.5283	18.0257	21.99m	29.9048	29.2359	1.64h
GCN (Joshi et al., 2019a)	SL, G	17.0141	58.7272	59.11s	29.7173	79.6063	6.67m	48.6151	110.2900	28.52m
GCN (Joshi et al., 2019a)	SL, BS	16.1878	51.0185	4.63m	30.3702	83.5523	38.02m	51.2593	121.7278	51.67m
GCN (Joshi et al., 2019a)	SL, BS*	16.2081	51.2079	3.97m	30.4258	83.8883	30.62m	51.0992	121.0357	3.23h
Att-GCRN(Fu et al., 2021)	SL+RL MCTS	10.7358	0.1563	20.62s + 1.33m	16.7471	1.2169	31.17s + 3.33m	23.5153	1.7179	43.94s + 6.68m
<b>UTSP (Ours)</b>	UL, Search	10.7289	<b>0.0918</b>	0.56m + 1.11m	16.6846	<b>0.8394</b>	1.37m + 1.33m	23.3903	<b>1.1770</b>	3.35m+ 2.67m

values are the results of the rounding problem. The total runtime of our method remains competitive w.r.t. all other learning baselines. On larger instances with  $n = 200, 500$  and  $1,000$ , we notice that traditional solvers and heuristics (Concorde, Gurobi and LKH3) fail to generate the optimal solutions within reasonable time when the size of problems grows. For RL/SL baselines, they generate results far away from ideal solutions, particularly for cases with  $n = 1,000$ . Our UTSP method is able to obtain **0.0918%**, **0.8394%** and **1.1770%** on TSP 200, 500 and 1,000, respectively. We remark that the UTSP takes a shorter total running time (inference + search) and outperform the existing learning baselines on these large instances. The gap between running time becomes more pronounced when the size increases to 1,000. More discussion between (Fu et al., 2021) and UTSP

can be found in Appendix 7.2.

Our model also takes less training time because we require very few training instances. Taking TSP 100 as an example, RL/SL needs 1 million training instances, and the total training time can take one day using a NVIDIA V100 GPU, while our method only takes about 40 minutes with 2,000 training instances. The training data size does not increase w.r.t. TSP size. Our training data consists of 2,000 instances for TSP 200, 500 and 1,000. At the same time, the UTSP model also remains very lightweight. On TSP 100, we use a 2-layer SAG with 64 hidden units and the model consists of 44,392 trainable parameters. In contrast, RL method in (Kool et al., 2019) takes approximately 700,000 parameters and the SL method in (Joshi et al., 2022) takes




Figure 6. The heat map  $\mathcal{H}$  generated using GCN on TSP 100. The diagonal elements are set to 0.  $X$ -axis and  $y$ -axis are the city indices.




Figure 7. The heat map  $\mathcal{H}$  generated using SAG on TSP 100. The diagonal elements are set to 0.  $X$ -axis and  $y$ -axis are the city indices.

approximately 350,000 parameters.

Our results indicate the UTSP algorithm is able to generate better solutions within a reasonable time. Our UL pipeline also generalizes well to unseen examples without requiring a large number of training samples. This is because the loss function in Equation 2 is fully differentiable w.r.t the parameters in SAG and we are able to train the model in an end-to-end fashion. In other words, given a heat map  $\mathcal{H}$ , the model learns to assign large weights to more promising edges and small weights to less promising ones through backpropagation without any prior knowledge of the ground truth or any exploration step. However, when using SL, the model learns from the TSP solutions, which fails when multiple solutions exist or the solutions are not optimal. While for RL, the model often encounters an exploration dilemma and is not guaranteed to converge (Bengio et al., 2021)(Joshi et al., 2019b). Overall, UTSP requires fewer training samples and has better generalization comparing to SL/RL models.

### 4.3. Expressive Power of GNNs

We aim to generate a non-smooth transition matrix  $\mathbb{T}$  and build an expressive heat map  $\mathcal{H}$  to guide the search algorithm. However, most GNNs aggregate information from adjacent nodes and these aggregation steps usually consist of local averaging operations, which can be interpreted as a low-pass filter and causes the oversmoothing problem (Wenkel et al., 2022). The low-pass model generates a smooth transition matrix  $\mathbb{T}$ , which finally makes the elements  $\mathcal{H}$  become indistinguishable. So it becomes difficult

to discriminate whether the edges belong to the optimal solution or not. In our model, we assume all nodes in the graph are connected, so every node has  $n - 1$  neighboring nodes. This means every node receives messages from all other nodes and we have a global averaging operation over the graph, this can lead to severe oversmoothing issue.

To avoid oversmoothing, one solution is to use shallow GNNs. However, this would result in narrow receptive fields and create the problem of underreaching (Barceló et al., 2020). In our model, we use SAG because this scattering-based method helps overcome the oversmoothing problem by combining band-pass wavelet filters with GCN-type filters (Min et al., 2022). Figure 5 illustrates the training loss on TSP 100 and the differences between our SAG model and the graph convolutional network (GCN) (Kipf & Welling, 2016), where GCN only performs low-pass filtering on graph signals (Nt & Maehara, 2019). When using GCN, the training loss decreases slowly, and the validation loss reaches a plateau after we train the model for 20 epochs. This is because the low-pass model generates a smooth  $\mathbb{T}$ . Such a smooth  $\mathbb{T}$  results in an indistinguishable  $\mathcal{H}$ , which harms the training process. Instead, we observe lower training and validation loss when using SAG; this suggests that SAG generates a more expressive representation which facilitates the training process.

Figure 6 and Figure 7 illustrate the generated heat maps using GCN and SAG on a TSP 100 instance, we choose this instance from the validation set randomly. When using the GCN, due to the oversmoothing problem, the model generates a smooth representation and  $\mathcal{H}$  becomes indistinguishable. The elements in  $\mathcal{H}$  have a small variance and




Figure 8. Average edge overlap coefficient  $\eta$  w.r.t. training epochs using SAG and GCN on TSP 100 ( $M = 10$ ).

most of them are  $\sim 0.01$ . Instead, the SAG generates a discriminative representation and the elements in the heat map have a larger variance.

Here, we train both GCN and SAG with the same loss function. So the differences illustrated in Figure 6 and 7 are the direct result of overcoming the oversmoothing problem.

## 5. Search Space Reduction

To understand what happens during our training process, we study how the prediction edge set  $\Pi$  changes with training time. As mentioned, let  $\Pi$  denote undirected edge set in  $\mathcal{H}'$ , and let  $\Gamma$  denote the ground truth edge set,  $\eta = |\Gamma \cap \Pi|/|\Gamma|$  is the extent of how good our prediction set  $\Pi$  covers the solution  $\Gamma$ . If  $\eta = 1$ , then  $\Gamma$  is a subset of  $\Pi$ , which means our prediction edge set successfully covers all ground truth edges. Similarly,  $\eta = 0.95$  means we cover 95% ground truth edges.

Figure 8 shows how the average overlap coefficient  $\eta$  changes with training epochs. We calculate the coefficient based on 1,000 validation instances in TSP 100. We notice that the coefficient quickly increases to  $\sim 98\%$  after we train SAG for 10 epochs. This suggests that the surrogate loss successfully encourages the SAG to put more weights on the more promising edges. We also compare the performance with GCN. Since the loss does not decrease significantly during our training when using GCN (shown in Figure 5), it is not surprising to see the average overlap coefficient of GCN always maintains at a relatively low level. After training the model for 100 epochs, SAG model has an average coefficient of 99.756% while GCN only has 33.893%.

Overall, the unsupervised learning training reduces the search space from 4950 edges to 583.134 edges with over 99% overlap accuracy. This helps explain why our search algorithm is able to perform well within reasonable time.

We then study how many cases where our prediction edge




Figure 9. Number of fully covered instances w.r.t. training epochs using SAG and GCN on TSP 100. The validation set consists of 1,000 samples ( $M = 10$ ).

set  $\Pi$  covers the ground truth solution. Figure 9 illustrates how the number of fully covered instances ( $\eta = 1$ ) changes with time. After training the model for 100 epochs, we observed 780 fully covered instances in 1,000 validation samples using SAG while 0 instances using GCN. Finally, we calculate the average of size  $|\Pi|$ . Our results show that SAG has an average size of 583.134 edges, while for GCN, the number is 738.739.

These results also indicate that there is a correspondence between the loss and the quality of our prediction. In most SL tasks such as classification or regression tasks, a smaller validation loss usually means we achieve better performance and the minimum of the loss corresponds to the global optimal solution (100% accuracy). However, it is no theoretical guarantee that our loss in Equation 2 is also a measure of the solution quality. Our empirical results demonstrate that a lower surrogate loss encourages the model to assign larger weights on the promising edges and reduces the search space. This implies that we can assess the quality of the generated heat maps using our loss in Equation 2.

We also compare the prediction edge sets and our results demonstrate that smooth representations fail to reduce the search space. Figure 10 and Figure 11 illustrate the difference of prediction edge sets between GCN and SAG. Figure 10 and Figure 11 are generated using the heat map in Figure 6 and Figure 7 with  $M = 10$ , respectively. The light green regions correspond to the prediction edge set  $\Pi$ . The  $x$ -axis and  $y$ -axis are the city indices, a light green box with position  $(i, j)$  means edge  $(i, j)$  belongs to  $\Pi$ . When using GCN, as shown in Figure 10, we observe more continuous light green regions comparing to Figure 11. As mentioned before, a low-pass model will enforce similarity on neighboring nodes and lead to unfavorable representations. The continuous regions in Figure 10 are the direct result of oversmoothness. We observe fewer continuous light green regions when using SAG, this suggest that the






Figure 10. Illustration of how the prediction edge set covers the ground truth edges using GCN. Light green: the prediction edge set  $\Pi$ ; blue:  $\Pi$  contains the ground truth edge, red:  $\Pi$  misses the ground truth edge.




Figure 11. Illustration of how the prediction edge set covers the ground truth edges using SAG. Light green: the prediction edge set  $\Pi$ ; blue:  $\Pi$  contains the ground truth edge, red:  $\Pi$  misses the ground truth edge.

model helps alleviate the oversmoothing problem and generates a more distinguishable representation. We further study how the prediction sets cover the ground truth solution. In Figure 10 and 11, blue and red boxes are the edges in ground truth solution. A blue box with position  $(i, j)$  corresponds to the condition that our prediction set  $\Pi$  covers the right edge  $(i, j)$ , while a red box at position  $(i, j)$  means there is a ground truth edge  $(i, j)$  but  $\Pi$  fails to cover it. When using GCN, we observe 118 red boxes and 82 blue boxes, this means the GCN misses 59 correct edges, while SAG’s prediction set successfully covers all the right edges.

Overall, the GCN’s prediction set has 875 edges (1,750 light green boxes) and SAG’s prediction set  $\Pi$  has 614 edges (1,228 light green boxes). Although the low-pass model has a larger prediction set  $\Pi$ , it still falls short of covering the right edges. This emphasizes the importance of including band-pass filters and overcoming the oversmoothness problem.

## 6. Conclusion

In this paper, we propose UTSP, an unsupervised learning method to solve the TSP. We build a surrogate loss that encourages the GNN to find the shortest path and satisfy the constraint that the path should be a Hamiltonian Cycle. The surrogate loss function does not rely on any labeled ground

truth solution and helps alleviate sparse reward problems in RL. UTSP uses a two-phase strategy. We first build a heat map based on the GNN’s output. The heat map is then fed into a search algorithm. Compared with RL/SL, our method vastly reduces training cost and takes fewer training samples. We further show that our UL training helps reduce the search space. This helps explain why the generated heat maps can guide the search algorithm. On the model side, our results indicate that a low-pass GNN will produce an indistinguishable representation due to the oversmoothing issue, which results in unfavorable heat maps and fails to reduce the search space. Instead, after incorporating band-pass operators into GNN, we can build efficient heat maps that successfully reduce search space. Our findings show that the expressive power of GNNs is critical for generating a non-smooth representation that helps find the solution.

In conclusion, UTSP is competitive with or outperforms other learning-based TSP heuristics in terms of solution quality and running speed. In addition, UTSP takes  $\sim 10\%$  of the number of parameters and  $\sim 0.2\%$  of (unlabeled) training samples, compared with reinforcement learning or supervised learning methods. Our UTSP framework demonstrates that, providing a surrogate loss and an expressive GNN, we can learn to reduce the search space and build a heuristic by exploiting a small amount of unlabeled data without any prior knowledge. These learned “algorithmic

priors” help facilitate the local search and lead to better solutions. Future direction includes designing more expressive GNNs (such as adding edge features) and using different surrogate loss functions. We anticipate that these concepts will extend to more (graph) combinatorial problems.

## References

- Applegate, D., Bixby, R., Chvatal, V., and Cook, W. Concorde tsp solver, 2006.
- Barceló, P., Kostylev, E. V., Monet, M., Pérez, J., Reutter, J., and Silva, J.-P. The logical expressiveness of graph neural networks. In *8th International Conference on Learning Representations (ICLR 2020)*, 2020.
- Bengio, Y., Lodi, A., and Prouvost, A. Machine learning for combinatorial optimization: a methodological tour d’horizon. *European Journal of Operational Research*, 290(2):405–421, 2021.
- Bresina, J. L. Heuristic-biased stochastic sampling. In *AAAI/IAAI, Vol. 1*, pp. 271–278, 1996.
- Croes, G. A. A method for solving traveling-salesman problems. *Operations research*, 6(6):791–812, 1958.
- Deudon, M., Cournut, P., Lacoste, A., Adulyasak, Y., and Rousseau, L.-M. Learning heuristics for the tsp by policy gradient. In *International Conference on the Integration of Constraint Programming, Artificial Intelligence, and Operations Research*, pp. 170–181. Springer, 2018.
- Fu, Z.-H., Qiu, K.-B., and Zha, H. Generalize a small pre-trained model to arbitrarily large tsp instances. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(8):7474–7482, 2021.
- Gomes, C. P., Selman, B., Kautz, H., et al. Boosting combinatorial search through randomization. *AAAI/IAAI*, 98: 431–437, 1998.
- Gomes, C. P., Selman, B., Crato, N., and Kautz, H. Heavy-tailed phenomena in satisfiability and constraint satisfaction problems. *Journal of automated reasoning*, 24(1): 67–100, 2000.
- Helsgaun, K. An effective implementation of the lin-kernighan traveling salesman heuristic. *European journal of operational research*, 126(1):106–130, 2000.
- Helsgaun, K. An extension of the lin-kernighan-helsgaun tsp solver for constrained traveling salesman and vehicle routing problems. *Roskilde: Roskilde University*, pp. 24–50, 2017.
- Joshi, C. K., Laurent, T., and Bresson, X. An efficient graph convolutional network technique for the travelling salesman problem. *arXiv preprint arXiv:1906.01227*, 2019a.
- Joshi, C. K., Laurent, T., and Bresson, X. On learning paradigms for the travelling salesman problem. *arXiv preprint arXiv:1910.07210*, 2019b.
- Joshi, C. K., Cappart, Q., Rousseau, L.-M., and Laurent, T. Learning the travelling salesperson problem requires rethinking generalization. *Constraints*, pp. 1–29, 2022.
- Kingma, D. P. and Ba, J. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- Kipf, T. N. and Welling, M. Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907*, 2016.
- Kool, W., van Hoof, H., and Welling, M. Attention, learn to solve routing problems! In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=ByxBFsRqYm>.
- Min, Y., Wenkel, F., Perlmutter, M., and Wolf, G. Can hybrid geometric scattering networks help solve the maximum clique problem? *arXiv preprint arXiv:2206.01506*, 2022.
- Nt, H. and Maehara, T. Revisiting graph neural networks: All we have is low-pass filters. *arXiv preprint arXiv:1905.09550*, 2019.
- Qiu, R., Sun, Z., and Yang, Y. Dimes: A differentiable meta solver for combinatorial optimization problems. *arXiv preprint arXiv:2210.04123*, 2022.
- Sylvester, J. J. *The Collected Mathematical Papers of James Joseph Sylvester...*, volume 3. University Press, 1909.
- Wenkel, F., Min, Y., Hirn, M., Perlmutter, M., and Wolf, G. Overcoming oversmoothness in graph convolutional networks via hybrid scattering networks. *arXiv preprint arXiv:2201.08932*, 2022.

## 7. Appendix

### 7.1. Training and Search Details

We train our model using Adam (Kingma & Ba, 2014). All models are trained using Nvidia V100 GPU. All the search-related parameters are listed in Table 3.  $M$  refers to the size of the candidate set of each city.  $K$  is the maximal number of edges we can remove in one action, and for each round of local search, we randomly select one number from the listed interval.  $T$  is the total number of actions we will try to expand one node. Here we keep the  $\alpha = 0$  to show that the quality of our unsupervised learned heat map is high. Lower  $\alpha$  means the local search algorithm focuses more on the edges with higher heat map value. Actually, in the experiments, we find the results are similar with  $\alpha \leq 1$ .

	$\alpha$	$\beta$	$M$	$K$	$T$
TSP-20	0	10	8	10	60
TSP-50	0	10	8	[5, 15)	150
TSP-100	0	10	8	[5, 35)	300
TSP-200	0	10	8	[10, 90)	600
TSP-500	0	50	5	[30, 130)	1000
TSP-1000	0	50	5	[10, 110)	2000

Table 3. Search parameters for all the TSP experiments.

### 7.2. Running Time Discussion

As discussed in (Kool et al., 2019), running time is important but hard to compare since it is affected by many factors. We report the clock time for solving all the test instances in Table 1. For the UTSP (our method) and the state-of-the-art learning-based method Att-GCRN (Fu et al., 2021), we run the search algorithm on exactly the same environment (one Intel Xeon Gold 6326) for a fair comparison. And for other baselines, we directly refer to the results from (Fu et al., 2021). So the time there are only for indicative purpose since the computing hardware is not the same.